

Scaling vs normalization

May 20, 2022

```
[1]: import pandas as pd
import numpy as np
from sklearn import datasets
wine = datasets.load_wine()
wine = pd.DataFrame(
    data=np.c_[wine['data'], wine['target']],
    columns=wine['feature_names'] + ['target']
)
```

```
[2]: wine
```

```
[2]:
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	\
0	14.23	1.71	2.43	15.6	127.0	2.80	
1	13.20	1.78	2.14	11.2	100.0	2.65	
2	13.16	2.36	2.67	18.6	101.0	2.80	
3	14.37	1.95	2.50	16.8	113.0	3.85	
4	13.24	2.59	2.87	21.0	118.0	2.80	
..	
173	13.71	5.65	2.45	20.5	95.0	1.68	
174	13.40	3.91	2.48	23.0	102.0	1.80	
175	13.27	4.28	2.26	20.0	120.0	1.59	
176	13.17	2.59	2.37	20.0	120.0	1.65	
177	14.13	4.10	2.74	24.5	96.0	2.05	

	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	\
0	3.06	0.28	2.29	5.64	1.04	
1	2.76	0.26	1.28	4.38	1.05	
2	3.24	0.30	2.81	5.68	1.03	
3	3.49	0.24	2.18	7.80	0.86	
4	2.69	0.39	1.82	4.32	1.04	
..	
173	0.61	0.52	1.06	7.70	0.64	
174	0.75	0.43	1.41	7.30	0.70	
175	0.69	0.43	1.35	10.20	0.59	
176	0.68	0.53	1.46	9.30	0.60	
177	0.76	0.56	1.35	9.20	0.61	

	od280/od315_of_diluted_wines	proline	target
0	3.92	1065.0	0.0
1	3.40	1050.0	0.0
2	3.17	1185.0	0.0
3	3.45	1480.0	0.0
4	2.93	735.0	0.0
..
173	1.74	740.0	2.0
174	1.56	750.0	2.0
175	1.56	835.0	2.0
176	1.62	840.0	2.0
177	1.60	560.0	2.0

[178 rows x 14 columns]

```
[3]: wine[['magnesium', 'total_phenols', 'color_intensity']].describe()
```

```
[3]:
```

	magnesium	total_phenols	color_intensity
count	178.000000	178.000000	178.000000
mean	99.741573	2.295112	5.058090
std	14.282484	0.625851	2.318286
min	70.000000	0.980000	1.280000
25%	88.000000	1.742500	3.220000
50%	98.000000	2.355000	4.690000
75%	107.000000	2.800000	6.200000
max	162.000000	3.880000	13.000000

As you can see all tree columns have different data distributions. So it may present a problem with distance based models like knn

```
[4]: wine.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 178 entries, 0 to 177
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   alcohol                              178 non-null    float64
1   malic_acid                           178 non-null    float64
2   ash                                  178 non-null    float64
3   alcalinity_of_ash                    178 non-null    float64
4   magnesium                            178 non-null    float64
5   total_phenols                        178 non-null    float64
6   flavanoids                           178 non-null    float64
7   nonflavanoid_phenols                 178 non-null    float64
8   proanthocyanins                      178 non-null    float64
9   color_intensity                      178 non-null    float64
10  hue                                  178 non-null    float64
```

```

11  od280/od315_of_diluted_wines  178 non-null    float64
12  proline                        178 non-null    float64
13  target                        178 non-null    float64
dtypes: float64(14)
memory usage: 19.6 KB

```

1 Let's try a KNN

```
[5]: from sklearn.neighbors import KNeighborsClassifier
```

```
[6]: knn_clf = KNeighborsClassifier()
```

```
[7]: # Let's split data first
from sklearn.model_selection import train_test_split
y = wine["target"]
X = wine.drop(["target"], axis=1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
↳ random_state=42)
```

```
[8]: model = knn_clf.fit(X_train, y_train)
```

```
[9]: preds = model.predict(X_test)
```

```
[10]: from sklearn.metrics import classification_report
print(classification_report(y_test, preds))
```

	precision	recall	f1-score	support
0.0	0.85	0.85	0.85	20
1.0	0.67	0.67	0.67	24
2.0	0.47	0.47	0.47	15
accuracy			0.68	59
macro avg	0.66	0.66	0.66	59
weighted avg	0.68	0.68	0.68	59

```
[11]: # current data distribution is
y_test.value_counts() / sum(y_test.value_counts())

# this enables us to compare with random model
```

```
[11]: 1.0    0.406780
0.0    0.338983
2.0    0.254237
Name: target, dtype: float64
```

2 Let's do the same thing but with data normalization

```
[12]: from sklearn.preprocessing import StandardScaler
      # create the scaler
      ss = StandardScaler()
```

```
[13]: X_train_norm = ss.fit_transform(X_train)
      X_test_norm = ss.fit_transform(X_test)
```

```
[14]: pd.DataFrame(X_train_norm).describe()
      # as can be seen now the values have average close to zero and std=1
```

```
[14]:
```

	0	1	2	3	4	\
count	1.190000e+02	1.190000e+02	1.190000e+02	1.190000e+02	1.190000e+02	
mean	9.143013e-16	-9.358177e-16	2.928563e-15	1.026257e-16	-2.549315e-16	
std	1.004228e+00	1.004228e+00	1.004228e+00	1.004228e+00	1.004228e+00	
min	-2.287878e+00	-1.380891e+00	-3.788813e+00	-2.570844e+00	-2.043016e+00	
25%	-8.027088e-01	-6.965013e-01	-5.132388e-01	-6.092008e-01	-8.512565e-01	
50%	-1.297567e-02	-4.653499e-01	-4.529965e-02	-2.213963e-02	-1.891681e-01	
75%	8.710539e-01	6.858750e-01	5.723800e-01	5.506030e-01	4.729203e-01	
max	2.191205e+00	3.069907e+00	3.211557e+00	2.984759e+00	4.048198e+00	

	5	6	7	8	9	\
count	1.190000e+02	1.190000e+02	1.190000e+02	1.190000e+02	1.190000e+02	
mean	2.287153e-15	-3.883448e-17	-4.688127e-16	6.260165e-16	5.392512e-16	
std	1.004228e+00	1.004228e+00	1.004228e+00	1.004228e+00	1.004228e+00	
min	-1.966736e+00	-1.648405e+00	-1.847646e+00	-2.028179e+00	-1.496884e+00	
25%	-8.912050e-01	-8.559681e-01	-7.423755e-01	-5.965104e-01	-8.179582e-01	
50%	3.284258e-02	1.235043e-01	-1.897403e-01	-4.987326e-02	-1.025798e-01	
75%	8.129811e-01	8.421121e-01	6.786864e-01	6.008852e-01	5.193893e-01	
max	2.426277e+00	3.017623e+00	2.336592e+00	3.455546e+00	2.631350e+00	

	10	11	12
count	1.190000e+02	1.190000e+02	1.190000e+02
mean	1.708251e-15	2.164468e-16	-1.016927e-16
std	1.004228e+00	1.004228e+00	1.004228e+00
min	-2.018458e+00	-1.809669e+00	-1.513264e+00
25%	-7.896365e-01	-1.050959e+00	-7.967631e-01
50%	1.545374e-02	2.545691e-01	-2.663237e-01
75%	6.934245e-01	7.877167e-01	7.162131e-01
max	3.193442e+00	1.922364e+00	2.629059e+00

```
[15]: model = knn_clf.fit(X_train_norm, y_train)
      preds = model.predict(X_test_norm)
      from sklearn.metrics import classification_report
      print(classification_report(y_test, preds))
```

```
# as can be seen in the results below, the results dramatically improved
```

	precision	recall	f1-score	support
0.0	0.95	1.00	0.98	20
1.0	1.00	0.88	0.93	24
2.0	0.88	1.00	0.94	15
accuracy			0.95	59
macro avg	0.94	0.96	0.95	59
weighted avg	0.95	0.95	0.95	59

3 Let's do the same thing but with min max scaler

```
[16]: from sklearn.preprocessing import MinMaxScaler
```

```
[20]: # create the scaler
minmax = MinMaxScaler()
X_train_mm = minmax.fit_transform(X_train)
X_test_mm = minmax.fit_transform(X_test)
```

```
[21]: pd.DataFrame(X_train_mm).describe()
```

```
[21]:
```

	0	1	2	3	4	5	\
count	119.000000	119.000000	119.000000	119.000000	119.000000	119.000000	
mean	0.510792	0.310257	0.541230	0.462748	0.335404	0.447696	
std	0.224204	0.225629	0.143454	0.180760	0.164865	0.228597	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.331579	0.153768	0.467914	0.353093	0.195652	0.244828	
50%	0.507895	0.205703	0.534759	0.458763	0.304348	0.455172	
75%	0.705263	0.464358	0.622995	0.561856	0.413043	0.632759	
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	

	6	7	8	9	10	11	\
count	119.000000	119.000000	119.000000	119.000000	119.000000	119.000000	
mean	0.353278	0.441573	0.369854	0.362597	0.387279	0.484902	
std	0.215221	0.240003	0.183129	0.243259	0.192680	0.269083	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.169831	0.264151	0.261076	0.164459	0.235772	0.203297	
50%	0.379747	0.396226	0.360759	0.337748	0.390244	0.553114	
75%	0.533755	0.603774	0.479430	0.488411	0.520325	0.695971	
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	

	12
count	119.000000

```

mean      0.365318
std       0.242431
min       0.000000
25%      0.172971
50%      0.301024
75%      0.538219
max       1.000000

```

```

[22]: model = knn_clf.fit(X_train_mm, y_train)
      preds = model.predict(X_test_mm)
      from sklearn.metrics import classification_report
      print(classification_report(y_test, preds))

      # as can be seen in the results below, the results dramatically improved

```

	precision	recall	f1-score	support
0.0	0.91	1.00	0.95	20
1.0	1.00	0.92	0.96	24
2.0	1.00	1.00	1.00	15
accuracy			0.97	59
macro avg	0.97	0.97	0.97	59
weighted avg	0.97	0.97	0.97	59

Again the results improved relative to the baseline, but they are equal or marginally better compared to standard normalization

```
[ ]:
```