

# Unsupervised Learning

May 23, 2022

```
[1]: # load data
import pandas as pd
```

```
[2]: df = pd.read_csv("cars.csv")
```

```
[3]: df
```

```
[3]:      mpg  cylinders  cubicinches  hp  weightlbs  time-to-60  year  \
0    14.0           8          350  165     4209         12   1972
1    31.9           4           89   71     1925         14   1980
2    17.0           8          302  140     3449         11   1971
3    15.0           8          400  150     3761         10   1971
4    30.5           4           98   63     2051         17   1978
..    ...           ...           ...  ...     ...         ...   ...
256  17.0           8          305  130     3840         15   1980
257  36.1           4           91   60     1800         16   1979
258  22.0           6          232  112     2835         15   1983
259  18.0           6          232  100     3288         16   1972
260  22.0           6          250  105     3353         15   1977
```

```
      brand
0      US.
1  Europe.
2      US.
3      US.
4      US.
..      ...
256     US.
257  Japan.
258     US.
259     US.
260     US.
```

[261 rows x 8 columns]

```
[4]: df.describe()
```

```
[4]:
```

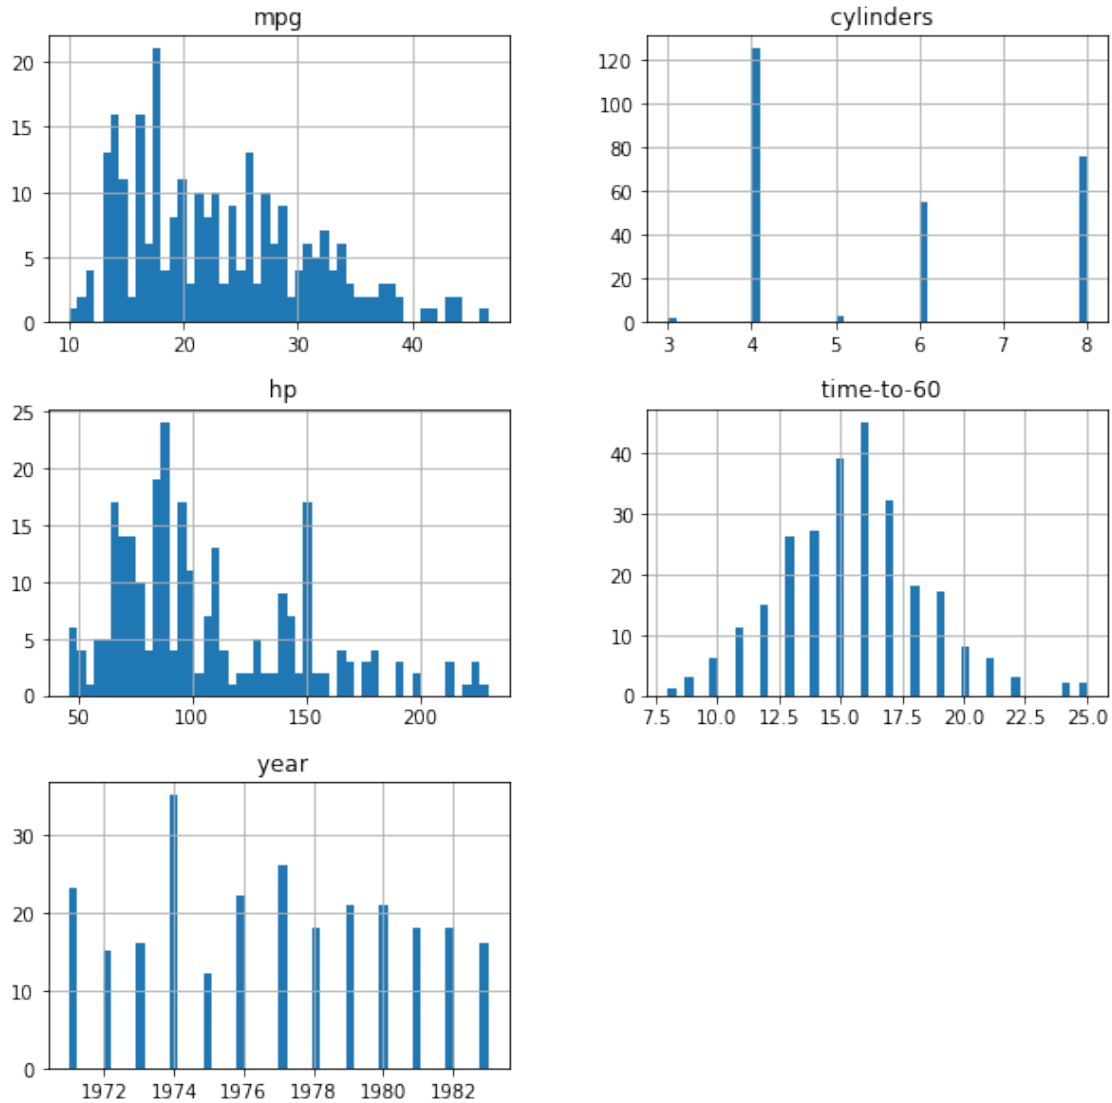
|       | mpg        | cylinders  | hp         | time-to-60 | year        |
|-------|------------|------------|------------|------------|-------------|
| count | 261.000000 | 261.000000 | 261.000000 | 261.000000 | 261.000000  |
| mean  | 23.144828  | 5.590038   | 106.360153 | 15.547893  | 1976.819923 |
| std   | 7.823570   | 1.733310   | 40.499959  | 2.910625   | 3.637696    |
| min   | 10.000000  | 3.000000   | 46.000000  | 8.000000   | 1971.000000 |
| 25%   | 16.900000  | 4.000000   | 75.000000  | 14.000000  | 1974.000000 |
| 50%   | 22.000000  | 6.000000   | 95.000000  | 16.000000  | 1977.000000 |
| 75%   | 28.800000  | 8.000000   | 138.000000 | 17.000000  | 1980.000000 |
| max   | 46.600000  | 8.000000   | 230.000000 | 25.000000  | 1983.000000 |

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 261 entries, 0 to 260
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              261 non-null    float64
1   cylinders        261 non-null    int64
2   cubicinches      261 non-null    object
3   hp              261 non-null    int64
4   weightlbs       261 non-null    object
5   time-to-60      261 non-null    int64
6   year            261 non-null    int64
7   brand           261 non-null    object
dtypes: float64(1), int64(4), object(3)
memory usage: 16.4+ KB
```

## 1 EDA

```
[6]: _ = df.hist(figsize=(10,10), bins=50)
```



From the above results, there is no need to fill data.

## 2 Normalize dataset

```
[7]: # fix column names
df.columns = list(map(lambda x: x.strip(" "), df.columns))
df.columns
```

```
[7]: Index(['mpg', 'cylinders', 'cubicinches', 'hp', 'weightlbs', 'time-to-60',
          'year', 'brand'],
          dtype='object')
```

```
[8]: # Bin age of the car
df["year_bin"] = pd.cut(df["year"], 5)
df["year_bin"].value_counts()
```

```
[8]: (1975.8, 1978.2]      66
     (1970.988, 1973.4]   54
     (1980.6, 1983.0]     52
     (1973.4, 1975.8]     47
     (1978.2, 1980.6]     42
     Name: year_bin, dtype: int64
```

```
[9]: from sklearn.preprocessing import LabelEncoder
```

```
[10]: le = LabelEncoder()
df["year"] = le.fit_transform(df["year_bin"])
df.drop("year_bin", axis=1, inplace=True)
```

```
[11]: set(df.brand)
```

```
[11]: {' Europe.', ' Japan.', ' US.'}
```

```
[12]: le = LabelEncoder()
df["brand"] = le.fit_transform(df["brand"])
```

```
[13]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
df = sc.fit_transform(df)
```

```
-----
ValueError                                Traceback (most recent call last)
Input In [13], in <cell line: 3>()
      1 from sklearn.preprocessing import StandardScaler
      2 sc = StandardScaler()
----> 3 df = sc.fit_transform(df)

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/sklearn/base.py:867, in
↳ TransformerMixin.fit_transform(self, X, y, **fit_params)
      863 # non-optimized default implementation; override when a better
      864 # method is possible for a given clustering algorithm
      865 if y is None:
      866     # fit method of arity 1 (unsupervised transformation)
--> 867     return self.fit(X, **fit_params).transform(X)
      868 else:
      869     # fit method of arity 2 (supervised transformation)
      870     return self.fit(X, y, **fit_params).transform(X)
```

```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/sklearn/preprocessing/
↳_data.py:809, in StandardScaler.fit(self, X, y, sample_weight)
    807 # Reset internal state before fitting
    808 self._reset()
--> 809 return self.partial_fit(X, y, sample_weight)

```

```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/sklearn/preprocessing/
↳_data.py:844, in StandardScaler.partial_fit(self, X, y, sample_weight)
    812 """Online computation of mean and std on X for later scaling.
    813
    814 All of X is processed as a single batch. This is intended for cases
    (...)
    841     Fitted scaler.
    842 """
    843 first_call = not hasattr(self, "n_samples_seen_")
--> 844 X = self._validate_data(
    845     X,
    846     accept_sparse=("csr", "csc"),
    847     dtype=FLOAT_DTYPES,
    848     force_all_finite="allow-nan",
    849     reset=first_call,
    850 )
    851 n_features = X.shape[1]
    853 if sample_weight is not None:

```

```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/sklearn/base.py:577, in
↳BaseEstimator._validate_data(self, X, y, reset, validate_separately,
↳**check_params)
    575     raise ValueError("Validation should be done on X, y or both.")
    576 elif not no_val_X and no_val_y:
--> 577     X = check_array(X, input_name="X", **check_params)
    578     out = X
    579 elif no_val_X and not no_val_y:

```

```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/sklearn/utils/validation
↳py:856, in check_array(array, accept_sparse, accept_large_sparse, dtype,
↳order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
↳ensure_min_features, estimator, input_name)
    854     array = array.astype(dtype, casting="unsafe", copy=False)
    855     else:
--> 856     array = np.asarray(array, order=order, dtype=dtype)
    857 except ComplexWarning as complex_warning:
    858     raise ValueError(
    859         "Complex data not supported\n{}\n".format(array)
    860     ) from complex_warning

```

```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/core/generic.py:
↳2064, in NDFrame.__array__(self, dtype)

```

```

2063 def __array__(self, dtype: npt.DTypeLike | None = None) -> np.ndarray:
-> 2064     return np.asarray(self._values, dtype=dtype)

```

```

ValueError: could not convert string to float: ''

```

```
[14]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 261 entries, 0 to 260
Data columns (total 8 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   mpg             261 non-null   float64
 1   cylinders       261 non-null   int64
 2   cubicinches     261 non-null   object
 3   hp              261 non-null   int64
 4   weightlbs      261 non-null   object
 5   time-to-60     261 non-null   int64
 6   year           261 non-null   int64
 7   brand          261 non-null   int64
dtypes: float64(1), int64(5), object(2)
memory usage: 16.4+ KB

```

```
[15]: # cubicinches and weightlbs also may have some strings in the data
```

```

[16]: import numpy as np
def find_empty_string(data):
    if data == '' or data == ' ':
        return np.nan
    else:
        return data

```

```

[17]: df.cubicinches = df.cubicinches.apply(find_empty_string)
df.weightlbs = df.weightlbs.apply(find_empty_string)

```

```
[18]: df.weightlbs.apply(find_empty_string).isna().sum()
```

```
[18]: 3
```

```
[19]: df.cubicinches.apply(find_empty_string).isna().sum()
```

```
[19]: 2
```

```
[24]: df.cubicinches.astype(int)
```

```

-----
ValueError

```

```

Traceback (most recent call last)

```

Input In [24], in <cell line: 1>()

```
----> 1 df.cubicinches.astype(int)
```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/core/generic.py:

```
→5912, in NDFrame.astype(self, dtype, copy, errors)
    5905     results = [
    5906         self.iloc[:, i].astype(dtype, copy=copy)
    5907         for i in range(len(self.columns))
    5908     ]
    5910 else:
    5911     # else, only a single dtype is given
-> 5912     new_data = self._mgr.astype(dtype=dtype, copy=copy, errors=errors)
    5913     return self._constructor(new_data).__finalize__(self,
→method="astype")
    5915 # GH 33113: handle empty frame or series
```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/core/internals/

```
→managers.py:419, in BaseBlockManager.astype(self, dtype, copy, errors)
    418 def astype(self: T, dtype, copy: bool = False, errors: str = "raise") -> T:
-> 419     return self.apply("astype", dtype=dtype, copy=copy, errors=errors)
```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/core/internals/

```
→managers.py:304, in BaseBlockManager.apply(self, f, align_keys,
→ignore_failures, **kwargs)
    302     applied = b.apply(f, **kwargs)
    303     else:
-> 304     applied = getattr(b, f)(**kwargs)
    305 except (TypeError, NotImplementedError):
    306     if not ignore_failures:
```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/core/internals/

```
→blocks.py:580, in Block.astype(self, dtype, copy, errors)
    562 """
    563 Coerce to the new dtype.
    564
    565 (...)
    576 Block
    577 """
    578 values = self.values
-> 580 new_values = astype_array_safe(values, dtype, copy=copy, errors=errors)
    582 new_values = maybe_coerce_values(new_values)
    583 newb = self.make_block(new_values)
```

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/core/dtypes/cast.

```
→py:1292, in astype_array_safe(values, dtype, copy, errors)
    1289     dtype = dtype.numpy_dtype
    1291 try:
```

```

-> 1292     new_values = astype_array(values, dtype, copy=copy)
1293 except (ValueError, TypeError):
1294     # e.g. astype_nansafe can fail on object-dtype of strings
1295     # trying to convert to float
1296     if errors == "ignore":

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/core/dtypes/cast.
py:1237, in astype_array(values, dtype, copy)
1234     values = values.astype(dtype, copy=copy)
1236 else:
-> 1237     values = astype_nansafe(values, dtype, copy=copy)
1239 # in pandas we don't store numpy str dtypes, so convert to object
1240 if isinstance(dtype, np.dtype) and issubclass(values.dtype.type, str):

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/core/dtypes/cast.
py:1154, in astype_nansafe(arr, dtype, copy, skipna)
1150 elif is_object_dtype(arr.dtype):
1151
1152     # work around NumPy brokenness, #1987
1153     if np.issubdtype(dtype.type, np.integer):
-> 1154         return lib.astype_intsafe(arr, dtype)
1156     # if we have a datetime/timedelta array of objects
1157     # then coerce to a proper dtype and recall astype_nansafe
1159     elif is_datetime64_dtype(dtype):

File ~/anaconda3/envs/EDIT/lib/python3.9/site-packages/pandas/_libs/lib.pyx:668
in pandas._libs.lib.astype_intsafe()

ValueError: cannot convert float NaN to integer

```

```

[25]: data = []
for i in df.cubicinches:
    try:
        data.append(int(i))
    except:
        print(i)

```

nan  
nan

```

[26]: cylinder_average = sum(data)/len(data)

```

```

[27]: df.cubicinches = df.cubicinches.fillna(cylinder_average).astype(int)

```

```

[28]: data = []
for i in df.weightlbs:
    try:

```



```

        data.append(int(i))
    except:
        print(i)

```

```

nan
nan
nan

```

```
[29]: weightlbs_average = sum(data)/len(data)
```

```
[30]: weightlbs_average
```

```
[30]: 3009.8333333333335
```

```
[31]: df.weightlbs = df.weightlbs.fillna(weightlbs_average).astype(int)
```

```
[32]: df
```

```
[32]:
```

|     | mpg  | cylinders | cubicinches | hp  | weightlbs | time-to-60 | year | brand |
|-----|------|-----------|-------------|-----|-----------|------------|------|-------|
| 0   | 14.0 | 8         | 350         | 165 | 4209      | 12         | 0    | 2     |
| 1   | 31.9 | 4         | 89          | 71  | 1925      | 14         | 3    | 0     |
| 2   | 17.0 | 8         | 302         | 140 | 3449      | 11         | 0    | 2     |
| 3   | 15.0 | 8         | 400         | 150 | 3761      | 10         | 0    | 2     |
| 4   | 30.5 | 4         | 98          | 63  | 2051      | 17         | 2    | 2     |
| ..  | ...  | ...       | ...         | ... | ...       | ...        | ...  | ...   |
| 256 | 17.0 | 8         | 305         | 130 | 3840      | 15         | 3    | 2     |
| 257 | 36.1 | 4         | 91          | 60  | 1800      | 16         | 3    | 1     |
| 258 | 22.0 | 6         | 232         | 112 | 2835      | 15         | 4    | 2     |
| 259 | 18.0 | 6         | 232         | 100 | 3288      | 16         | 0    | 2     |
| 260 | 22.0 | 6         | 250         | 105 | 3353      | 15         | 2    | 2     |

```
[261 rows x 8 columns]
```

```
[33]: sc = StandardScaler()
df = sc.fit_transform(df)
```

### 3 Apply TSNE

```
[34]: from sklearn.manifold import TSNE
```

```
[35]: # We want to get TSNE embedding with 2 dimensions
n_components = 2
tsne = TSNE(n_components)
tsne_result = tsne.fit_transform(df)
tsne_result.shape
```

```

/home/local/FARFETCH/tiago.cabo/anaconda3/envs/EDIT/lib/python3.9/site-
packages/sklearn/manifold/_t_sne.py:795: FutureWarning: The default
initialization in TSNE will change from 'random' to 'pca' in 1.2.
  warnings.warn(
/home/local/FARFETCH/tiago.cabo/anaconda3/envs/EDIT/lib/python3.9/site-
packages/sklearn/manifold/_t_sne.py:805: FutureWarning: The default learning
rate in TSNE will change from 200.0 to 'auto' in 1.2.
  warnings.warn(

```

[35]: (261, 2)


```

[36]: import seaborn as sns
import matplotlib.pyplot as plt
tsne_result_df = pd.DataFrame({'tsne_1': tsne_result[:,0], 'tsne_2':
    ↪tsne_result[:,1]})
fig, ax = plt.subplots(1)
sns.scatterplot(x='tsne_1', y='tsne_2', data=tsne_result_df, ax=ax,s=120)
lim = (tsne_result.min()-5, tsne_result.max()+5)
ax.set_xlim(lim)
ax.set_ylim(lim)
ax.set_aspect('equal')
ax.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.0)

```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

[36]: <matplotlib.legend.Legend at 0x7f6a1c700970>



From the T-SNE plot, we can see that 3 or 4 clusters may exist

## 4 Let's Apply k-means

```
[37]: from sklearn.cluster import KMeans
```

```
[38]: # Let's just confirm that the data is normal distributed
pd.DataFrame(df).describe()
```

```
[38]:
```

|       | 0             | 1             | 2             | 3             | 4 \           |
|-------|---------------|---------------|---------------|---------------|---------------|
| count | 2.610000e+02  | 2.610000e+02  | 2.610000e+02  | 2.610000e+02  | 2.610000e+02  |
| mean  | 3.143505e-16  | 2.124737e-16  | 8.443650e-17  | 9.831429e-17  | -2.475670e-16 |
| std   | 1.001921e+00  | 1.001921e+00  | 1.001921e+00  | 1.001921e+00  | 1.001921e+00  |
| min   | -1.683385e+00 | -1.497144e+00 | -1.223544e+00 | -1.493239e+00 | -1.648229e+00 |
| 25%   | -7.997404e-01 | -9.191048e-01 | -9.197567e-01 | -7.758132e-01 | -8.918596e-01 |
| 50%   | -1.466117e-01 | 2.369740e-01  | -4.134443e-01 | -2.810368e-01 | -1.248703e-01 |
| 75%   | 7.242265e-01  | 1.393053e+00  | 9.305848e-01  | 7.827325e-01  | 7.719173e-01  |
| max   | 3.003774e+00  | 1.393053e+00  | 2.339054e+00  | 3.058704e+00  | 2.344835e+00  |

|       | 5             | 6             | 7             |
|-------|---------------|---------------|---------------|
| count | 2.610000e+02  | 2.610000e+02  | 2.610000e+02  |
| mean  | -2.043916e-16 | 2.911145e-17  | 1.795073e-16  |
| std   | 1.001921e+00  | 1.001921e+00  | 1.001921e+00  |
| min   | -2.598203e+00 | -1.402393e+00 | -1.833878e+00 |
| 25%   | -5.328295e-01 | -6.888950e-01 | -5.574989e-01 |
| 50%   | 1.556284e-01  | 2.460339e-02  | 7.188801e-01  |
| 75%   | 4.998573e-01  | 7.381017e-01  | 7.188801e-01  |
| max   | 3.253689e+00  | 1.451600e+00  | 7.188801e-01  |


```
[39]: inertias = []
mapping2 = {}
K = range(1, 10)
for k in K:
    # Building and fitting the model
    kmeanModel = KMeans(n_clusters=k).fit(df)
    kmeanModel.fit(df)

    inertias.append(kmeanModel.inertia_)

    mapping2[k] = kmeanModel.inertia_
```

```
[40]: plt.plot(K, inertias, 'bx-')
plt.xlabel('Values of K')
plt.ylabel('Inertia')
```

```
plt.title('The Elbow Method using Inertia')
plt.show()
```




```
[41]: # plot data with 3 clusters
kmeanModel = KMeans(n_clusters=3).fit(df)
kmeanModel.fit(df)
```

```
[41]: KMeans(n_clusters=3)
```

```
[42]: labels = kmeanModel.labels_
```

```
[43]: import seaborn as sns
import matplotlib.pyplot as plt
tsne_result_df = pd.DataFrame({'tsne_1': tsne_result[:,0], 'tsne_2':
    ↪tsne_result[:,1], 'label':labels})
fig, ax = plt.subplots(1)
sns.scatterplot(x='tsne_1', y='tsne_2', hue='label', data=tsne_result_df,
    ↪ax=ax, s=120)
lim = (tsne_result.min()-5, tsne_result.max()+5)
ax.set_xlim(lim)
ax.set_ylim(lim)
ax.set_aspect('equal')
ax.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.0)
```

[43]: <matplotlib.legend.Legend at 0x7f6a1c34c7c0>



```
[44]: # Let's try with 4 clusters
# plot data with 3 clusters
kmeanModel = KMeans(n_clusters=4).fit(df)
kmeanModel.fit(df)

# store results
labels = kmeanModel.labels_

tsne_result_df = pd.DataFrame({'tsne_1': tsne_result[:,0], 'tsne_2':
    ↪tsne_result[:,1], 'label':labels})
fig, ax = plt.subplots(1)
sns.scatterplot(x='tsne_1', y='tsne_2',hue='label', data=tsne_result_df,
    ↪ax=ax,s=120)
lim = (tsne_result.min()-5, tsne_result.max()+5)
ax.set_xlim(lim)
ax.set_ylim(lim)
ax.set_aspect('equal')
ax.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.0)
```

[44]: <matplotlib.legend.Legend at 0x7f6a1c27e130>

